

Heuristics for generating abstract test cases from \mathcal{R} ebeca model

Pröstur Thorarensen (throstur11@ru.is)

April 16, 2015

Abstract

Software testing can be a difficult task for humans and is a prime candidate for automation. Concurrent systems are no exception, even when employing paradigms such as the Actor model. Model checking goes a long way to verifying software correctness, but models always abstract away from the implementation. Mutant testing is a promising method to pinpoint programming errors, by combining the abstract model of a system with mutants of the implementation, it is possible to automatically generate an extensive suite of test cases with ease. The research demonstrates that software testing can easily be supplemented by automation and that even the simplistic methods employed in this study are capable of producing reliable results.

Contents

| | | |
|----------|--|-----------|
| 1 | Introduction | 3 |
| 2 | Contribution | 3 |
| 3 | Related Work | 4 |
| 4 | Method | 4 |
| 4.1 | Modeling the System Under Test | 4 |
| 4.2 | Coverage Criteria | 4 |
| 4.3 | Generating Test Cases | 4 |
| 4.3.1 | Abstract Test Cases | 4 |
| 4.3.2 | Concrete Test Cases | 5 |
| 4.4 | Mutant Generation | 5 |
| 4.5 | Test Case Evaluation | 5 |
| 5 | Results | 6 |
| 6 | Discussion | 6 |
| 7 | Conclusion | 6 |
| 8 | Future Work | 7 |
| | Appendices | 8 |
| A | Rebeca Model - Full Code | 8 |
| B | Implementation - Full Code | 15 |

1 Introduction

A model describing a System Under Test (SUT) is usually an abstract, partial presentation of the SUT's desired behavior. Model-based testing is using such a model of the SUT to generate abstract test cases and then mapping those abstract test cases to executable test cases based on the back-end code. Test cases derived from such a model are at the same level of abstraction as the model. An abstract test case cannot be directly executed against a SUT and an executable test suite needs to be derived from a corresponding abstract test suite.

Model checking is a method of formally verifying finite-state concurrent systems according to a given specification. The specification is generally expressed as temporal logic formulas which are verified either with explicit-state model checking or symbolic model checking. Symbolic model checking is more efficient states can be represented as sets of states rather than as single states.

Rebeca is an actor-based modeling language designed in an effort to bridge the gap between formal verification approaches and real applications. Rebeca introduces the ability to analyze a group of reactive objects as single components of a system in the actor model.

Since we cannot test software with all inputs, **coverage criteria** are used to decide which test inputs to use. The software testing community believes that effective use of coverage criteria makes it more likely that test engineers will find faults in a program and provides informal assurance that the software is of high quality and reliability.

Mutation testing is used to design new software tests and evaluate the quality of existing software tests. It involves creating *mutants* modified versions of the SUT that are based on *mutation operators*. Each mutant $m \in M$ where M is the set of mutants for a given artifact will lead to a test requirement¹ and yields more test requirements than any other test criterion in most situations (provided that the mutation operators are well-designed). In practice, if software contains a fault, there will usually be a set of mutants that can only be killed by a test case that also detects that fault [Ammann and Offutt, 2008].

2 Contribution

Coverage criteria is necessary when testing all but the most trivial software. Coverage criteria can however be difficult to decide upon for complicated systems. Software developers and testers are often faced with the task of testing software. This can be error-prone and time consuming, making testing a useful target for automation.

In this study, we modeled a case study using Rebeca and employed model checking to ascertain that the model is faithful to the specification. Following this, we generated the state space of the model and extracted (random) traces through the graph. Although employing some heuristics for which traces are extracted would have been nice, this has been left as future work. The extracted traces were transformed into test cases for the SUT.

We used Offutt et al. to generate mutants for the SUT. Each mutant was evaluated by the test cases in an attempt to find mutants that could not be killed. These mutants act as the coverage criteria for the randomly generated traces.

¹Criterion 5.32 on Mutation Coverage

3 Related Work

TODO!

4 Method

4.1 Modeling the System Under Test

The case study intended as the SUT is the NASA GMSEC Message Bus, the specification of which was provided by Fraunhofer CESE. The model of the system was implemented as a `reactiveclass` in *Rebeca*. Simply modeling the SUT as a `reactiveclass` was not sufficient for our purposes - the system is reactive and as such requires inputs from external entities to avoid idling. For this reason, a `reactiveclass` modeling a user of the SUT was implemented, non-deterministically sending messages to the SUT. To avoid a combinatorial explosion in generating the state space, the amount of messages sent must have a reasonably low upper bound - 20 was selected for this study.

It should be noted that an actual implementation of the SUT was not used in this study. Instead, the model was re-implemented as a Java class following the specification of the system². This was done for a few reasons, notably giving control over used frameworks to the researchers and to double-check that the model is faithful to the specification. The SUT implementation was implemented as a Akka Java actor, as the Akka Testkit framework provides excellent tools to test actor based systems.

4.2 Coverage Criteria

The coverage criteria for the test suite generated in this research is the set of all mutants that must be killed. Recall that *Criterion 5.32* in Ammann and Offut [2008] states that each mutant $m \in M$ where M is the set of mutants for a given artifact will lead to a test requirement - each test requirement produced by the elements in M is a part of our coverage criteria. The test suite consists of a tuple (M, T) , where T is a list³ of traces. Only messages to and from the SUT are considered (all other transitions in the state space are considered as τ transitions and were discarded from the test case generation). The state space of the model was exported from Afra (*Rebeca* IDE) as an `.aut`⁴ file suitable for visualization with software such as CADP.

4.3 Generating Test Cases

4.3.1 Abstract Test Cases

Random traces were extracted from the state space. The traces act as abstract test cases for the SUT, as the behavior of an actor based model to an outside observer can be fully modeled based only on the

²Using this approach proved to be quite useful. We were able to pinpoint an error in the model due to incorrect handling of an edge case which was correct in the implementation (which our properties for model checking did not cover). Although a complete implementation should also catch this, using the complete implementation increases the complexity of the project and reduces the amount of time that could be dedicated solely to the research topic at hand.

³It would be more useful to have a *set* of traces, but since randomness does not guarantee uniqueness, we will make do with a list.

⁴ALDEBARAN/*AUTomaton* file format. Specification: <http://www.inrialpes.fr/vasy/cadp/man/aut.html#sect2>

messages that are being sent between actors.

In order to extract individual traces from the *Rebeca* model, we implemented simple Java software to create a graph of the state space based on an `.aut` file as input, with the output being a list of traces through the graph. Branching decisions in the state space were made at random. The traces extracted from the state space are then used as abstract test cases. The code can be found in the `trace-extract` package of the `trace-mutants` GitHub repository.

4.3.2 Concrete Test Cases

The `trace-extract` package also includes a Java class for transforming the abstract test cases to concrete test cases. The `TestGenerator` class generates a Scala test spec for the Akka Testkit framework⁵.

4.4 Mutant Generation

Mutants were generated using the muJava Offutt et al. [2015] mutation system for Java programs. Some slight modifications were made to the source code of muJava to facilitate mutant compilation. The artifacts of the modifications are available in the `gen-mutants` package of the `trace-mutants` GitHub repository. Every available mutant operator was applied to the original source, generating a total of 316 mutants.

Class mutants can be generated by muJava by using class-level mutation operators. This was not done since the simplified implementation consisted mostly of a single class (with inner classes as message types). muJava defines 12 method-level operators which were all applied to the SUT implementation. Arithmetic, Conditional and Logical operators can be either *replaced*, *deleted*, or *inserted*⁶ in the code. Relational, shift and assignment operators can be replaced but insertion or deletion is not done. Ma and Offut [2005] provide further details for mutation operators.

4.5 Test Case Evaluation

The test cases can be evaluated with the `evaluate-mutants` package of the `trace-mutants` GitHub repository. `evaluate-mutants` uses the file system to set up individual environments for each mutant to be compiled and tested in and then uses Scala-SBT to execute the test cases.

The result of a test case evaluation can give different information depending on the results. If all mutants are killed, either:

- The coverage of the test suite is insufficient; or
- The SUT conforms to the model and the test suite satisfies the coverage criteria.

Otherwise:

- There is a bug in the SUT, representative of a mutant that was not killed; or

⁵The `TestGenerator` class hard-codes the messages that should be sent and received by the SUT as a result of fast prototyping, but the code is available on GitHub

⁶Only unary operators can be inserted.

- There is a discrepancy between the SUT and the model.

5 Results

Although muJava attempted to generate 356 mutants, only 316 mutants of our implementation of the SUT could be compiled. 277 of those mutants modified the `onReceive(Object message)` method whereas 79 modified other parts of the program, including the message constructors and the overloaded `boolean equals(Object other)` methods.

We found that 50 random traces from the state space did not generate a sufficiently diverse set of traces to reliably kill all mutants. In one run, 5 mutants survived the 50 randomly chosen traces⁷: COR_21, ODL_33, SDL_17, SDL_19, VDL_14. The SDL mutants modified the `boolean equals(Object other)` method whereas the other mutants modified the `onReceive(Object message)` method. ODL_33 and VDL_13 resulted in the same mutation (removing a predicate from an OR conditional) whereas COR_21 replaced the OR operator to an AND operator.

Similar results were encountered during development of the test framework. Since randomness plays such a large part in determining whether or not a set of traces through the state space graph provides adequate coverage. We found that generating 1000 traces resulted in all mutants being killed every time, although theoretically it is of course always possible to generate unhelpful test cases when using randomness. We also tested 250 traces⁸ with all mutants being killed. Due to the nature of randomness, we cannot guarantee that any number of traces will invariantly kill all mutants (provided that killing all mutants is at all possible) - at least not until some heuristics are applied to trace extraction.

We were able to find one discrepancy between the model and the implementation, in which an edge case was being handled incorrectly in the model but correctly in the implementation. This demonstrates that in order to kill all mutants, it is important that the model and implementation must be faithful to each other. Mutants generated from a correct implementation but an incorrect model are more difficult to kill and create caps in the coverage, just as mutants generated from an incorrect implementation but a correct model.

6 Discussion

TODO: mention SBT problems on parallel distributed file systems

7 Conclusion

Foo!

⁷data/0050_4_0.traces

⁸data/0250_0_0.traces

8 Future Work

Despite the contribution to the field made by this study, there is still much research left to be explored. Most notably, the work should be repeated with a complete implementation of a system rather than with an implementation that is quite similar to the model itself. Increased complexity in a system's implementation can introduce hidden side effects that could possibly affect the results of the test cases.

Furthermore, due to the unexpected issues with Scala-SBT on parallel distributed file systems, the build system (which acts also as the test runner) should be replaced by a more suitable alternative. This has been done in part, but remains to be tested⁹.

A highly desirable improvement on the current work would be to add some sort of heuristics to the selection of traces from the state space of the system. Currently the traces are chosen at random with no guarantees of any actual coverage - with actual heuristics in place it will be possible to guarantee some sort of coverage rather than depending on randomness.

References

Paul Ammann and Jeff Offut. *Introduction to software testing*. Cambridge University Press, New York, 2008. ISBN 978-0-521-88038-1.

Fraunhofer CESE. Fraunhofer center for experimental software engineering. <http://cs.gmu.edu/~offutt/mujava/>, 2015.

Yu-Seung Ma and Jeff Offut. Description of method-level mutation operators for java. 2005.

Jeff Offutt, Yu-Seung Ma, Yong Rae Kwon, and Nan Li. mujava. <http://cs.gmu.edu/~offutt/mujava/>, 2015.

Scala-SBT. Simple build tool. <http://www.scala-sbt.org/>, 2015.

Prøstur Thorarensen. Github repository for trace-mutants. <https://github.com/ThrosturX/trace-mutants>, 2015.

⁹The build system was replaced with a bash script in an effort to allow tests to be run on a high-performance computing cluster, but the cluster was ongoing maintenance and the required tools were out-of-date, if at all present.

Appendices

A Rebeca Model - Full Code

```
reactiveclass MessageBus(15) {
  knownrebecs {
    Attacker att;
  }

  statevars {
    // ATT vars
    boolean connectionExists;
    boolean connected;
    boolean subscribed;
    int subs;
    // callbacks and messages
    int messages;
    int numCallbackSubscriptions;
    // other
    int destroyedMsgs;
  }

  msgsrv initial() {
    connectionExists = false;
    connected = false;
    subscribed = false;
    subs = 0;
    messages = 0;
    numCallbackSubscriptions = 0;
    destroyedMsgs = 0;
  }

  msgsrv createConnection() {
    if (sender == att) {
      if (!connectionExists)
      {
        connectionExists = true;
        att.ack(1);
      }
      else
      {
        att.fail(); // error
      }
    }
  }

  msgsrv destroyConnection() {
    if (sender == att) {
      if (connectionExists)
      {
```



```

        connected = false; // disconnect
        numCallbackSubscriptions = 0; // delete subscriptions
        subscribed = false; // delete subscriptions
        destroyedMsgs += messages;
        messages = 0; // delete messages
        connectionExists = false; // destroy connection object
        att.ack(2);
    }
    else
    {
        att.fail(); // error
    }
}
}

msgsrv connect() {
    if (sender == att) {
        if (connectionExists && !connected)
        {
            connected = true;
            att.ack(3);
        }
        else
        {
            att.fail(); // error
        }
    }
}

msgsrv disconnect() {
    if (sender == att) {
        if (connected)
        {
            connected = false;
            att.ack(4);
        }
        else
        {
            att.fail(); // error
        }
    }
}

msgsrv subscribe() {
    if (sender == att) {
        if (subs < 2)
        {
            subs = subs + 1;
        }
        if (connected && !subscribed)
        {
            subscribed = true;

```

```

        att.ack(5);
    }
    else
    {
        att.fail(); // error
    }
}
}

msgsrv subscribeCallback() {
    if (sender == att) {
        if (connected && numCallbackSubscriptions < 10)
        {
            numCallbackSubscriptions += 1;
            att.ack(7);
        }
        else
        {
            att.fail(); // error
        }
    }
}

msgsrv unsubscribe() {
    if (sender == att) {
        if (subs >= 0)
        {
            subs = subs - 1;
        }
        if (connected && subscribed)
        {
            subscribed = false;
            att.ack(6);
        } else if (connected && numCallbackSubscriptions > 0) {
            if (subs < 2)
            {
                subs = subs + 1;
            }
            att.ack(6);
        }
        else
        {
            att.fail(); // error
        }

        if (connected) {
            // delete all callbacks unconditionally
            numCallbackSubscriptions = 0; // what if we are not connected?
        }
    }
}
}

```

```

msgsrv unsubscribeCallback() {
    if (sender == att) {
        if (connected && numCallbackSubscriptions > 0)
        {
            numCallbackSubscriptions -= 1;
            att.ack(8);
        }
        else
        {
            att.fail(); // error
        }
    }
}

msgsrv publish() {
    if (sender == att) {
        if (connected)
        {
            if (subscribed || numCallbackSubscriptions > 0) {
                messages += 1; // add the message to the queue

                int i; // invoke callback function for each subscription
                for (i = 0; i < numCallbackSubscriptions; i = i + 1)
                {
                    att.callback();
                }
            }
            att.ack(9);
        }
        else
        {
            att.fail(); // error
        }
    }
}

msgsrv getMessage() {
    if (sender == att) {
        if (connected && messages > 0)
        {
            att.receive();
            messages -= 1;
        }
        else
        {
            att.fail(); // error
        }
    }
}
}

// the Attacker is basically just a "dumb" app that doesn't follow the

```

```

    API properly
    // its purpose is to try to find counterexamples to the properties
    reactiveclass Attacker(32) {
        knownrebecs {
            MessageBus bus;
        }

        statevars {
            int connected; // 0 = disconnected; 1 = connected; 2 = connecting
            int rMsgs; // 'received'/processed messages
            int rCallbacks; // # of invoked callbacks
            int cSubs; // # of callback subscriptions (owned)
            int published; // # of published messages
            int acks;
            boolean stop;
            boolean hasUnsubscribed;
        }

        msgsrv initial() {
            connected = 0;
            rMsgs = 0;
            rCallbacks = 0;
            cSubs = 0;
            published = 0;
            acks = 0;
            stop = false;
            hasUnsubscribed = false;
            self.ack(0);
            self.crazy(-1); // continue after failures
        }

        msgsrv receive() {
            rMsgs += 1;
        }

        msgsrv callback() {
            rCallbacks += 1;
        }

        // this message server receives success codes and simulates a
        ridiculous test-run
        // this server only receives ACCEPT acknowledgements, the ID specifies
        what action succeeded
        msgsrv ack(int id) {
            acks = acks + 1;
            if (acks > 20) { // only simulate 20 "stupid" steps
                stop = true;
            }

            // handle changes
            if (id == 1) {
                connected = 2;
            }
        }
    }

```

```

} else if (id == 2 | id == 4) {
    connected = 0;
} else if (id == 3) {
    connected = 1;
}

if (id == 0) {          // used to activate command chain
    connected = 2;
    bus.createConnection();
} else if (id == 9) { // published
    published += 1;
    self.idle();
} else if (!stop) {
    int old = id;
    self.crazy(old);
} else if (stop && connected == 0 && (id == 2 || id == -1)) {
    self.idle();
} else if (stop) {
    hasUnsubscribed = true;
    bus.destroyConnection();
} else {
    self.idle();
}
}

// try something else
msgsrv fail() {
    acks = acks + 1;
    if (acks < 20) {
        self.crazy(-1);
    } else {
        self.idle();
    }
}

// performs a random action
msgsrv crazy(int old) {
    int id;

    id = ?(0, 1, 2, 3, 4, 5, 6, 7, 8, 9);

    if (id == 0) {
        bus.createConnection();
    } else if (id == 1) { // connection object created
        bus.connect();
    } else if (id == 2) {
        if (old == 2) {
            stop = true;
        }
    } else if (id == 3) {
        bus.subscribe();
    } else if (id == 4) {

```

```

        hasUnsubscribed = true; // destroying the connection loses
                               subscriptions
        bus.destroyConnection();
    } else if (id == 5) {
        bus.publish();
    } else if (id == 6) {
        bus.disconnect();
    } else if (id == 7) {
        bus.unsubscribeCallback();
    } else if (id == 8) {
        hasUnsubscribed = true;
        bus.unsubscribe();
    } else if (id == 9) {
        bus.subscribeCallback();
    }
}

msgsrv idle() {
    self.idle();
}

}

main {
    MessageBus bus(att):();
    Attacker att(bus):();
}

```

The *Rebeca* model includes two **rebecs**: the message bus itself and an application (**Attacker**) that sends messages to the message bus at random. This ensures that as long as the **MessageBus rebec** is correct, the generated state space will represent the state space of a correct implementation of the SUT.

B Implementation - Full Code

```
import akka.actor.UntypedActor;
import akka.event.Logging;
import akka.event.LoggingAdapter;

import java.util.Queue;
import java.util.LinkedList;

public class JBus extends UntypedActor {

    public interface BusMessage {
        // intentionally left blank
    }

    public final static class CreateConnection implements BusMessage {
        private final int id;

        public CreateConnection(int id) {
            this.id = id;
        }
    }

    public final static class DestroyConnection implements BusMessage {
        private final int id;

        public DestroyConnection(int id) {
            this.id = id;
        }
    }

    public final static class Connect implements BusMessage {
        private final int id;

        public Connect(int id) {
            this.id = id;
        }
    }

    public final static class Disconnect implements BusMessage {
        private final int id;

        public Disconnect(int id) {
            this.id = id;
        }
    }

    public final static class Subscribe implements BusMessage {
        private final int id;

        public Subscribe(int id) {
            this.id = id;
        }
    }
}
```

```

    }
}

public final static class SubscribeCallback implements BusMessage {
    private final int id;

    public SubscribeCallback(int id) {
        this.id = id;
    }
}

public final static class Unsubscribe implements BusMessage {
    private final int id;

    public Unsubscribe(int id) {
        this.id = id;
    }
}

public final static class UnsubscribeCallback implements BusMessage {
    private final int id;

    public UnsubscribeCallback(int id) {
        this.id = id;
    }
}

public final static class Publish implements BusMessage {
    private final int id;
    String message;

    public Publish(int id, String message) {
        this.id = id;
        this.message = message;
    }
}

public final static class GetMessage implements BusMessage {
    private final int id;

    public GetMessage(int id) {
        this.id = id;
    }
}

public final static class Success implements BusMessage {
    @Override
    public boolean equals(Object o) {
        if (o instanceof Success) {
            return true;
        }
        return false;
    }
}

```



```

    }
}

public final static class Error implements BusMessage {
    @Override
    public boolean equals(Object o) {
        if (o instanceof Error) {
            return true;
        }
        return false;
    }
}

private boolean connectionObjectExists = false;
private boolean connected = false;
private boolean subscribed = false;
private int callbacks = 0;
private int callbackInvocations = 0; // internal counter
private Queue<String> messages = new LinkedList<>();

public void onReceive(Object message) throws Exception {
    if (message instanceof CreateConnection) {
        if (!connectionObjectExists) {
            connectionObjectExists = true;
            getSender().tell(new Success(), getSelf());
        } else {
            getSender().tell(new Error(), getSelf());
        }
    } else if (message instanceof DestroyConnection) {
        if (connectionObjectExists) {
            connectionObjectExists = false;
            connected = false;
            subscribed = false;
            callbacks = 0;
            messages.clear();
            getSender().tell(new Success(), getSelf());
        } else {
            getSender().tell(new Error(), getSelf());
        }
    } else if (message instanceof Connect) {
        if (connectionObjectExists && !connected) {
            connected = true;
            getSender().tell(new Success(), getSelf());
        } else {
            getSender().tell(new Error(), getSelf());
        }
    } else if (message instanceof Disconnect
        && connected) {
        connected = false;
        getSender().tell(new Success(), getSelf());
    } else if (message instanceof Subscribe
        && connected && !subscribed) {

```

```

        subscribed = true;
        getSender().tell(new Success(), getSelf());
    } else if (message instanceof SubscribeCallback
        && connected && callbacks < 10) {
        callbacks += 1;
        getSender().tell(new Success(), getSelf());
    } else if (message instanceof Unsubscribe
        && connected && (subscribed || callbacks > 0)) {
        subscribed = false;
        callbacks = 0; // possibly differs from model
        getSender().tell(new Success(), getSelf());
    } else if (message instanceof UnsubscribeCallback
        && connected && (callbacks > 0)) {
        callbacks -= 1;
        getSender().tell(new Success(), getSelf());
    } else if (message instanceof Publish
        && connected) {
        /// BEWARE POSSIBLE BUG!
        if (subscribed || callbacks > 0) {
            messages.add(((Publish) message).message);
        }
        getSender().tell(new Success(), getSelf());
    } else if (message instanceof GetMessage
        && connected && messages.size() > 0) {
        String msg = messages.poll();
        getSender().tell(new Success(), getSelf());
    } else {
        /// unhandled(message);
        getSender().tell(new Error(), getSelf());
    }
}
}

```

The implementation is analogous to the *Rebeca* model. Details are omitted but the the implementation and model are faithful to each other.